```
from tensorflow.keras.datasets import mnist
# Load the MNTST dataset
(x train, y train), (x test, y test) = mnist.load data()
    Downloading data from <a href="https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz">https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz</a>
     11490434/11490434 — 1s 0us/step
# Check the shape of the data
print(f"x_train shape: {x_train.shape}")
print(f"y train shape: {y train.shape}")
print(f"x test shape: {x test.shape}")
print(f"v test shape: {v test.shape}")
→ x_train shape: (60000, 28, 28)
     v train shape: (60000,)
     x_test shape: (10000, 28, 28)
    v test shape: (10000,)
# Preprocess the data
# Flatten the 28x28 images to 1D arrays of 784 pixels
x train = x train.reshape(x train.shape[0], 784)
x_{\text{test}} = x_{\text{test.reshape}}(x_{\text{test.shape}}[0], 784)
Double-click (or enter) to edit
# Normalize pixel values to be between 0 and 1
x train = x train.astype('float32') / 255
x_test = x_test.astype('float32') / 255
Double-click (or enter) to edit
```

```
# One-hot encode the labels
v train = to categorical(v train, 10)
v test = to categorical(v test, 10)
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout
# Build the Feed-Forward Neural Network model
model = Sequential([
   # Input layer (784 features from flattened 28x28 images)
   Dense(512, activation='relu', input_shape=(784,)),
   Dropout(0.2), # Add dropout to prevent overfitting
   # Hidden layer
   Dense(256, activation='relu'),
   Dropout(0.2),
   # Another hidden layer
   Dense(128, activation='relu'),
   Dropout(0.2),
   # Output layer (10 classes for digits 0-9)
   Dense(10, activation='softmax')
])
    /usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input shape`,
      super().__init__(activity_regularizer=activity_regularizer, **kwargs)
from tensorflow.keras.optimizers import Adam
# Compile the model
model.compile(optimizer=Adam(learning_rate=0.0005), loss='categorical_crossentropy', metrics=['accuracy'])
```

Print model summary
model.summary()

→ Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401,920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32,896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1,290

Total params: 567,434 (2.16 MB)
Trainable params: 567,434 (2.16 MB)
Non-trainable params: 0 (0.00 B)

```
history = model.fit(
    x_train, y_train,
    batch_size=128,
    epochs=15, # Number of iterations
    validation_split=0.1,
    verbose=1
```

```
Epoch 5/15
                      ——— 11s 18ms/step - accuracy: 0.9812 - loss: 0.0602 - val accuracy: 0.9797 - val loss: 0.067
422/422 -
Epoch 6/15
                         — 10s 18ms/step — accuracy: 0.9846 — loss: 0.0497 — val accuracy: 0.9805 — val loss: 0.06!
422/422 -
Epoch 7/15
                      —— 10s 17ms/step - accuracy: 0.9864 - loss: 0.0433 - val accuracy: 0.9822 - val loss: 0.063
422/422 -
Epoch 8/15
                      ----- 7s 17ms/step - accuracy: 0.9881 - loss: 0.0367 - val accuracy: 0.9815 - val loss: 0.064!
422/422 -
Epoch 9/15
                         - 8s 18ms/step - accuracy: 0.9888 - loss: 0.0355 - val accuracy: 0.9820 - val loss: 0.067
422/422 -
Epoch 10/15
                       —— 7s 16ms/step - accuracy: 0.9903 - loss: 0.0286 - val accuracy: 0.9812 - val loss: 0.0653
422/422 ----
Epoch 11/15
                       —— 11s 18ms/step - accuracy: 0.9922 - loss: 0.0241 - val accuracy: 0.9813 - val loss: 0.069
422/422 -
Epoch 12/15
                        — 10s 18ms/step – accuracy: 0.9917 – loss: 0.0244 – val accuracy: 0.9847 – val loss: 0.063
422/422 ----
Epoch 13/15
422/422 ----
                        — 7s 17ms/step – accuracy: 0.9932 – loss: 0.0215 – val accuracy: 0.9827 – val loss: 0.0664
Epoch 14/15
422/422 ----
                       Epoch 15/15
422/422 ---
                      ——— 11s 18ms/step - accuracy: 0.9942 - loss: 0.0169 - val accuracy: 0.9833 - val loss: 0.072
```

```
# Evaluate the model on test data
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc:.4f}')
```

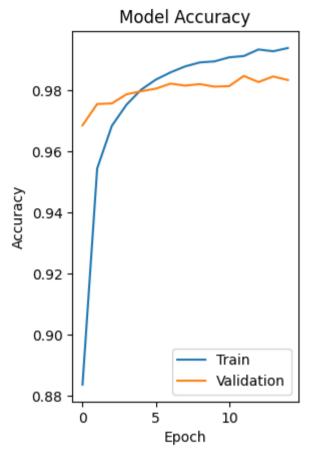
```
313/313 — 1s 3ms/step - accuracy: 0.9798 - loss: 0.0874 Test accuracy: 0.9827
```

```
import numpy as np

# Check model predictions distribution
predictions = model.predict(x_test)
predicted_classes = np.argmax(predictions, axis=1)
print("\nDistribution of predicted digits:")
```

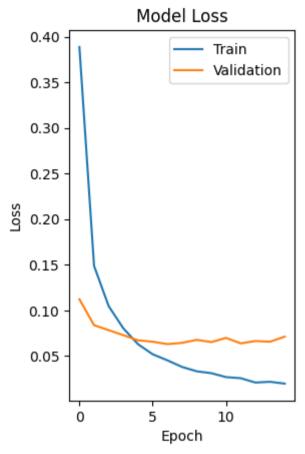
```
for i in range(10):
   print(f"Digit {i}: {np.sum(predicted classes == i)} predictions")
   313/313 — 1s 2ms/step
    Distribution of predicted digits:
    Digit 0: 978 predictions
    Digit 1: 1135 predictions
    Digit 2: 1052 predictions
    Digit 3: 1003 predictions
    Digit 4: 991 predictions
    Digit 5: 896 predictions
    Digit 6: 960 predictions
    Digit 7: 1019 predictions
    Digit 8: 965 predictions
    Digit 9: 1001 predictions
import matplotlib.pyplot as plt
# Plot training history
plt.figure(figsize=(12, 4))
→ <Figure size 1200x400 with 0 Axes>
    <Figure size 1200x400 with 0 Axes>
# Plot training & validation accuracy
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='lower right')
```

<matplotlib.legend.Legend at 0x79c1407edd90>



```
# Plot training & validation loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')
```





```
plt.tight_layout()
plt.show()
```

→ <Figure size 640x480 with 0 Axes>

Function to make predictions on new images
def predict_digit(image):

Predict the digit from an image.

```
Args:
        image: A 28x28 grayscale image array
    Returns:
        predicted digit (0-9)
    111111
    # Preprocess the image - flatten from 28x28 to 784 pixels
    image = image.reshape(1, 784)
    # Make prediction
    prediction = model.predict(image)
    print("prediction: ", prediction)
    # Get the predicted digit
    digit = prediction.argmax()
    confidence = prediction[0][digit]
    return digit, confidence
# Try with some test images
num samples = 5
plt.figure(figsize=(15, 3))
for i in range(num_samples):
   # Get a sample image
    image = x_test[i]
   # Make prediction
    digit, confidence = predict_digit(image)
   # Plot
    plt.subplot(1, num_samples, i+1)
    plt.imshow(image.reshape(28, 28)
    , cmap='gray')
    plt.title(f"Predicted: {digit}\nConfidence: {confidence:.2f}")
    plt.axis('off')
```

```
→ 1/1 — 0s 36ms/step
    prediction: [[2.0393348e-11 4.5532689e-09 3.9061856e-09 1.0482846e-08 2.7588144e-11
      4.6158702e-12 8.9156655e-13 9.9999988e-01 9.3772266e-13 7.3429852e-0811
    1/1 — 0s 37ms/step
    prediction: [[2.5952335e-10 2.7809580e-07 9.9999952e-01 2.1251353e-07 2.4667226e-12
      3.4523456e-10 3.1349763e-11 1.6199108e-09 2.0911288e-09 1.7063746e-1311
    1/1 — 0s 37ms/step
    prediction: [[9.6679276e-10 9.9999917e-01 1.2459350e-07 8.1800455e-10 3.4777266e-08
      5.5170544e-09 1.9610836e-08 5.2141951e-07 6.9251954e-08 2.9973843e-1011
    1/1 — 0s 36ms/step
    prediction: [[9.99965429e-01 1.65396727e-07 4.11489873e-06 7.34131831e-08
      5.45342743e-07 1.17587000e-07 2.76622559e-05 1.17647794e-07
      2.17447820e-07 1.48691890e-0611
    1/1 — 0s 36ms/step
    prediction: [[7.6735951e-09 2.9498544e-08 6.6759185e-09 1.5046782e-10 9.9999762e-01
      3.1631673e-09 6.0314008e-08 9.8671410e-07 2.8907570e-09 1.1459441e-06]]
         Predicted: 7
                                Predicted: 2
                                                      Predicted: 1
                                                                            Predicted: 0
                                                                                                  Predicted: 4
                                                                          Confidence: 1.00
        Confidence: 1.00
                              Confidence: 1.00
                                                    Confidence: 1.00
                                                                                                 Confidence: 1.00
plt.tight layout()
plt.show()
```

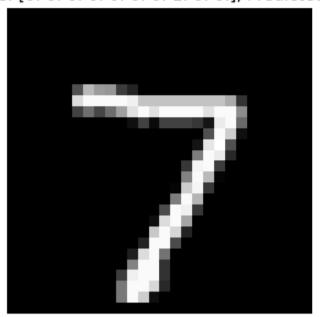
```
→ <Figure size 640x480 with 0 Axes>
```

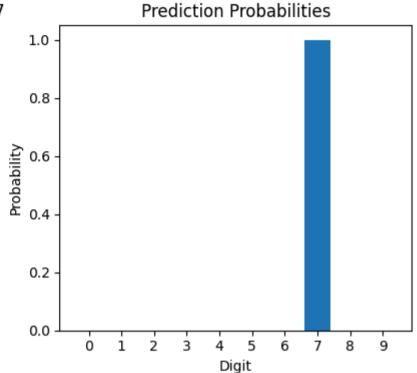
```
# Let's examine some predictions in detail
def display_prediction_details(index):
   """Show detailed prediction information for a given test example"""
   image = x test[index].reshape(1, 784)
   true_label = y_test[index]
```

```
# Get raw predictions (probabilities)
   pred probs = model.predict(image)[0]
   pred class = np.argmax(pred probs)
   # Display the image
   plt.figure(figsize=(8, 4))
   plt.subplot(1, 2, 1)
   plt.imshow(image.reshape(28, 28), cmap='gray')
   plt.title(f"True: {true_label}, Predicted: {pred class}")
   plt.axis('off')
   # Display the prediction probabilities
   plt.subplot(1, 2, 2)
   plt.bar(range(10), pred probs)
   plt.xticks(range(10))
   plt.xlabel('Digit')
   plt.ylabel('Probability')
   plt.title('Prediction Probabilities')
   plt.tight_layout()
   plt.show()
   print(f"Probabilities for each digit:")
   for i, prob in enumerate(pred probs):
        print(f"Digit {i}: {prob:.4f}")
# Check predictions for first few examples
for i in range(5):
   display_prediction_details(i)
```

_____ 0s 36ms/step

True: [0. 0. 0. 0. 0. 0. 1. 0. 0.], Predicted: 7





Probabilities for each digit:

Digit 0: 0.0000

Digit 1: 0.0000

Digit 2: 0.0000

Digit 3: 0.0000

Digit 4: 0.0000

Digit 5: 0.0000

Digit 6: 0.0000

Digit 7: 1.0000

Digit 8: 0.0000

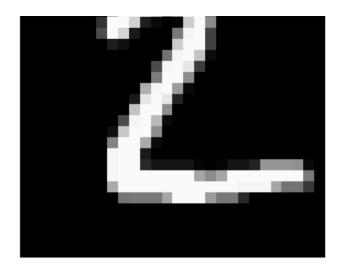
Digit 9: 0.0000 1/1 — 0s 37ms/step

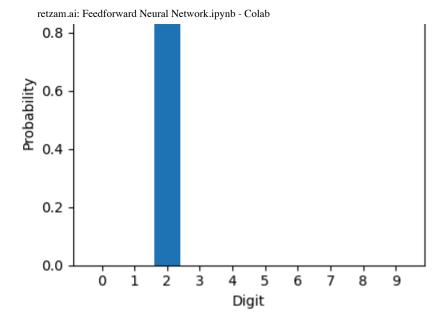
True: [0. 0. 1. 0. 0. 0. 0. 0. 0.], Predicted: 2



Prediction Probabilities

1.0



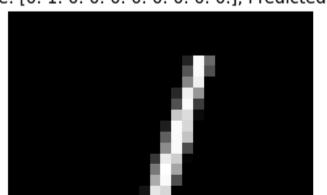


Probabilities for each digit:
Digit 0: 0.0000
Digit 1: 0.0000
Digit 2: 1.0000
Digit 3: 0.0000
Digit 4: 0.0000
Digit 5: 0.0000
Digit 6: 0.0000
Digit 7: 0.0000
Digit 8: 0.0000
Digit 9: 0.0000

1/1 -

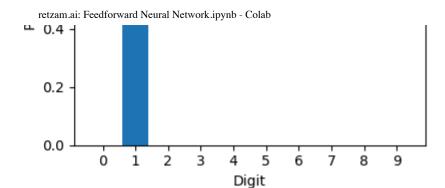
True: [0. 1. 0. 0. 0. 0. 0. 0. 0.], Predicted: 1

_____ 0s 37ms/step



Prediction Probabilities 1.0 0.8 \$\frac{\text{Ailing a probabilities}}{\text{Discrete of the probabilities}} = \frac{\text{Ailing a probabilities}}{\text{Discrete of the probabilities}} = \fr

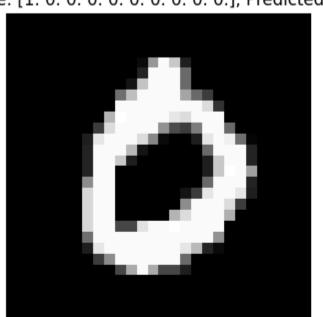




Probabilities for each digit: Digit 0: 0.0000 Digit 1: 1.0000 Digit 2: 0.0000 Digit 3: 0.0000 Digit 4: 0.0000 Digit 5: 0.0000 Digit 6: 0.0000 Digit 7: 0.0000 Digit 8: 0.0000 Digit 9: 0.0000 1/1 ----

_____ 0s 38ms/step

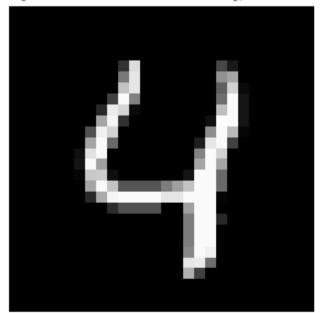
True: [1. 0. 0. 0. 0. 0. 0. 0. 0.], Predicted: 0



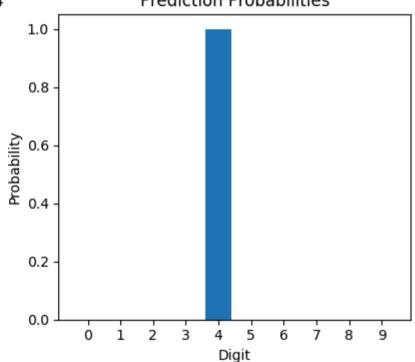
Prediction Probabilities 1.0 0.8 -Probability 0.2 -

```
Probabilities for each digit:
Digit 0: 1.0000
Digit 1: 0.0000
Digit 2: 0.0000
Digit 3: 0.0000
Digit 4: 0.0000
Digit 5: 0.0000
Digit 6: 0.0000
Digit 7: 0.0000
Digit 8: 0.0000
Digit 9: 0.0000
```

True: [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.], Predicted: 4



Prediction Probabilities



Probabilities for each digit:

Digit 0: 0.0000 Digit 1: 0.0000

Digit 2: 0.0000

Digit 3: 0.0000 Digit 4: 1.0000 Digit 5: 0.0000 Digit 6: 0.0000 Digit 7: 0.0000 Digit 8: 0.0000 Digit 9: 0.0000

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.